

Manaranjan Pradhan

manaranjan@enablecloud.com

This notebook is given as part of **Data Science for everyone** workshop.

(Forwarding this document to others is strictly prohibited.)

## Classification: Building multiple models ¶

- Decision Tree
- Random Forest
- Naive Bayes

In [2]:

```
import pandas as pd
import numpy as np
```

In [3]:

```
churn = pd.read_csv( "churn.csv" )
```

In [4]:

```
churn.head()
```

Out[4]:

	State	Account Length	Area Code	Phone	Int'l Plan	VMail Plan	VMail Message	Day Mins	Day Calls	Day Charge	..
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	..
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	..
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	..
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	..
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	..

5 rows × 21 columns



In [5]:

```
churn.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 3333 entries, 0 to 3332
Data columns (total 21 columns):
State                3333 non-null object
Account Length      3333 non-null int64
Area Code           3333 non-null int64
Phone               3333 non-null object
Int'l Plan          3333 non-null object
VMail Plan          3333 non-null object
VMail Message       3333 non-null int64
Day Mins            3333 non-null float64
Day Calls           3333 non-null int64
Day Charge          3333 non-null float64
Eve Mins            3333 non-null float64
Eve Calls           3333 non-null int64
Eve Charge          3333 non-null float64
Night Mins          3333 non-null float64
Night Calls         3333 non-null int64
Night Charge        3333 non-null float64
Intl Mins           3333 non-null float64
Intl Calls          3333 non-null int64
Intl Charge         3333 non-null float64
CustServ Calls      3333 non-null int64
Churn?              3333 non-null object
dtypes: float64(8), int64(8), object(5)
memory usage: 572.9+ KB
```

In [83]:

```
churn.columns
```

Out[83]:

```
Index(['State', 'Account Length', 'Area Code', 'Phone', 'Int'l Plan',
       'VMail Plan', 'VMail Message', 'Day Mins', 'Day Calls', 'Day C
       harge',
       'Eve Mins', 'Eve Calls', 'Eve Charge', 'Night Mins', 'Night Ca
       lls',
       'Night Charge', 'Intl Mins', 'Intl Calls', 'Intl Charge',
       'CustServ Calls', 'Churn?'],
      dtype='object')
```

## Data Cleaning

In [7]:

```
drop_columns = ['State', 'Area Code', 'Phone', 'Churn?',
                "Int'l Plan", "VMail Plan"]
```

In [8]:

```
churn_new = churn.drop( drop_columns, axis = 1 )
```

In [9]:

```
churn_new.head()
```

Out[9]:

	Account Length	VMail Message	Day Mins	Day Calls	Day Charge	Eve Mins	Eve Calls	Eve Charge	Night Mins	Night Calls
0	128	25	265.1	110	45.07	197.4	99	16.78	244.7	91
1	107	26	161.6	123	27.47	195.5	103	16.62	254.4	103
2	137	0	243.4	114	41.38	121.2	110	10.30	162.6	104
3	84	0	299.4	71	50.90	61.9	88	5.26	196.9	89
4	75	0	166.7	113	28.34	148.3	122	12.61	186.9	121

In [10]:

```
churn_new["Intl Plan"] = np.where( churn["Int'l Plan"] == 'yes',  
                                   1, 0 )
```

In [11]:

```
churn_new["VMail Plan"] = np.where( churn["Int'l Plan"] == 'yes',  
                                   1, 0 )
```

In [12]:

```
churn_new["churn"] = np.where( churn['Churn?'] == 'True.',1,0)
```

In [13]:

```
churn_new.head()
```

Out[13]:

	Account Length	VMail Message	Day Mins	Day Calls	Day Charge	Eve Mins	Eve Calls	Eve Charge	Night Mins	Night Calls
0	128	25	265.1	110	45.07	197.4	99	16.78	244.7	91
1	107	26	161.6	123	27.47	195.5	103	16.62	254.4	103
2	137	0	243.4	114	41.38	121.2	110	10.30	162.6	104
3	84	0	299.4	71	50.90	61.9	88	5.26	196.9	89
4	75	0	166.7	113	28.34	148.3	122	12.61	186.9	121

In [14]:

```
churn_new.columns = [name.replace(' ', '_') for name in churn_new.columns]
```

In [84]:

```
churn_new.head()
```

Out[84]:

	Account_Length	VMail_Message	Day_Mins	Day_Calls	Day_Charge	Eve_Charge
0	128	25	265.1	110	45.07	197.47
1	107	26	161.6	123	27.47	195.50
2	137	0	243.4	114	41.38	121.24
3	84	0	299.4	71	50.90	61.90
4	75	0	166.7	113	28.34	148.50



## Exploration

In [16]:


```
import matplotlib as plt
import seaborn as sn
import matplotlib.pyplot as pyplot
%matplotlib inline
```

In [17]:

```
pyplt.figure(figsize=(10, 8))  
sn.distplot( churn_new.Account_Length )
```

Out[17]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x921b4a8>




In [18]:

```
g = sns.FacetGrid( churn_new, col="churn", size = 10 )
g.map( sns.distplot, "Account_Length" )
```

Out[18]:

<seaborn.axisgrid.FacetGrid at 0x7ee3828>



In [19]:

```
churn_by_cs = pd.crosstab( churn_new.churn,
                             churn_new.CustServ_Calls )
```

In [20]:

```
churn_by_cs
```

Out[20]:

CustServ_Calls	0	1	2	3	4	5	6	7	8	9
churn										
0	605	1059	672	385	90	26	8	4	1	0
1	92	122	87	44	76	40	14	5	1	2

In [21]:

```
churn_by_cs = pd.DataFrame( churn_by_cs.unstack() ).reset_index()  
churn_by_cs.head()
```

Out[21]:

	CustServ_Calls	churn	0
0	0	0	605
1	0	1	92
2	1	0	1059
3	1	1	122
4	2	0	672

In [22]:

```
churn_by_cs = pd.crosstab( churn_new.churn,  
                           churn_new.CustServ_Calls )  
                           .apply( lambda calls: calls / calls.sum(), axis = 0 ) * 100  
churn_by_cs = pd.DataFrame( churn_by_cs.unstack() ).reset_index()  
churn_by_cs.head()
```

Out[22]:

	CustServ_Calls	churn	0
0	0	0	86.800574
1	0	1	13.199426
2	1	0	89.669771
3	1	1	10.330229
4	2	0	88.537549

In [23]:


```
churn_by_cs.columns = [ "CustServ_Calls", "churn", "total" ]
```

In [24]:

```
pyplt.figure(figsize=(10, 8))  
  
sn.barplot( churn_by_cs.CustServ_Calls,  
            churn_by_cs.total,  
            hue = churn_by_cs.churn )
```

Out[24]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x943cda0>



**The charges and minutes are highly correlated. So, only one variable can be used.**




In [25]:

```
churn_new.corr()
```

Out[25]:

	Account_Length	VMail_Message	Day_Mins	Day_Calls	Da
Account_Length	1.000000	-0.004628	0.006216	0.038470	0.0
VMail_Message	-0.004628	1.000000	0.000778	-0.009548	0.0
Day_Mins	0.006216	0.000778	1.000000	0.006750	1.0
Day_Calls	0.038470	-0.009548	0.006750	1.000000	0.0
Day_Charge	0.006214	0.000776	1.000000	0.006753	1.0
Eve_Mins	-0.006757	0.017562	0.007043	-0.021451	0.0
Eve_Calls	0.019260	-0.005864	0.015769	0.006462	0.0
Eve_Charge	-0.006745	0.017578	0.007029	-0.021449	0.0
Night_Mins	-0.008955	0.007681	0.004323	0.022938	0.0
Night_Calls	-0.013176	0.007123	0.022972	-0.019557	0.0
Night_Charge	-0.008960	0.007663	0.004300	0.022927	0.0
Intl_Mins	0.009514	0.002856	-0.010155	0.021565	-0.
Intl_Calls	0.020661	0.013957	0.008033	0.004574	0.0
Intl_Charge	0.009546	0.002884	-0.010092	0.021666	-0.
CustServ_Calls	-0.003796	-0.013263	-0.013423	-0.018942	-0.
Intl_Plan	0.024735	0.008745	0.049396	0.003755	0.0
VMail_Plan	0.024735	0.008745	0.049396	0.003755	0.0
churn	0.016541	-0.089728	0.205151	0.018459	0.2



In [26]:

```
import re
churn_cols = churn_new.columns
```

In [27]:

```
churn_drop_cols = [ x for x in churn_cols if
                    re.search( 'Charge', x ) or
                    re.search( 'Calls', x )]
```

In [28]:

```
churn_drop_cols
```

Out[28]:

```
['Day_Calls',  
 'Day_Charge',  
 'Eve_Calls',  
 'Eve_Charge',  
 'Night_Calls',  
 'Night_Charge',  
 'Intl_Calls',  
 'Intl_Charge',  
 'CustServ_Calls']
```

In [29]:

```
churn_drop_cols = churn_drop_cols[:-1]
```

In [30]:

```
churn_final = churn_new.drop( churn_drop_cols, axis = 1 )
```

In [31]:

```
churn_final.head()
```

Out[31]:

	Account_Length	VMail_Message	Day_Mins	Eve_Mins	Night_Mins	Intl_M
0	128	25	265.1	197.4	244.7	10.0
1	107	26	161.6	195.5	254.4	13.7
2	137	0	243.4	121.2	162.6	12.2
3	84	0	299.4	61.9	196.9	6.6
4	75	0	166.7	148.3	186.9	10.1

## Split Dataset into train and test

In [32]:

```
from sklearn.linear_model import LogisticRegression  
from sklearn.cross_validation import cross_val_score  
from sklearn.cross_validation import train_test_split
```

In [33]:

```
feature_cols = churn_final.columns[:-1].tolist()
```

In [49]:

```
X_train, X_test, y_train, y_test = train_test_split( churn_final[feature_cols],
                                                    churn_final.churn,
                                                    test_size = 0.3,
                                                    random_state = 21 )
```

## Building a Logistic Regression Model

In [35]:

```
logreg = LogisticRegression()
scores = cross_val_score(logreg, X_train, y_train, cv=10, scoring='accuracy')
```

In [36]:

```
scores
```

Out[36]:

```
array([ 0.85897436,  0.84615385,  0.82478632,  0.83333333,  0.8589743
6,
        0.8583691 ,  0.86266094,  0.87553648,  0.87931034,  0.875
])
```

In [37]:

```
scores.mean()
```

Out[37]:

```
0.85730990906549587
```

In [38]:

```
logreg.fit( X_train, y_train )
```

Out[38]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercep
t=True,
                    intercept_scaling=1, max_iter=100, multi_class='ovr',
                    penalty='l2', random_state=None, solver='liblinear', tol=0.
0001,
                    verbose=0)
```

In [39]:

```
y_predict = logreg.predict( X_test )
```

In [40]:


```
from sklearn import metrics
```

In [41]:

```
cm = metrics.confusion_matrix( y_test, y_predict )
```

In [42]:

```
sn.heatmap(cm, annot=True, fmt='.2f' );
```



In [43]:

```
metrics.accuracy_score( y_test, y_predict )
```

Out[43]:

```
0.8609999999999999
```

## Building a Decision Tree Model

In [44]:

```
from sklearn.tree import DecisionTreeClassifier
```

In [45]:

```
treeClf = DecisionTreeClassifier()  
scores = cross_val_score(treeClf, X_train, y_train, cv=10, scoring='accuracy')
```

In [46]:

```
scores
```

Out[46]:

```
array([ 0.87606838,  0.88461538,  0.86324786,  0.89316239,  0.8931623  
9,  
        0.90987124,  0.88841202,  0.89270386,  0.86637931,  0.8706896  
6])
```

## Verifying the optimal tree depth

In [50]:

```
max_depth_range = range(1, 12)

# list to store the average RMSE for each value of max_depth
accuracy_scores = []


# use LOOCV with each value of max_depth
for depth in max_depth_range:
    treereg = DecisionTreeClassifier(max_depth=depth, random_state=1)
    scores = cross_val_score(treereg, X_train, y_train, cv=14, scoring='accuracy')
    accuracy_scores.append( scores.mean() )
```

In [51]:

```
# plot max_depth (x-axis) versus RMSE (y-axis)
pyplt.plot( max_depth_range, accuracy_scores )
pyplt.xlabel( 'max_depth' )
pyplt.ylabel('Accuracy Scores')
```

Out[51]:

<matplotlib.text.Text at 0xc00e668>



In [78]:

```
# max_depth=3 was best, so fit a tree using that parameter
churn_tree_clf = DecisionTreeClassifier(max_depth=6, random_state=1)
churn_tree_clf.fit(X_train, y_train)
```

Out[78]:

```
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=
6,
                        max_features=None, max_leaf_nodes=None, min_samples_leaf=
1,
                        min_samples_split=2, min_weight_fraction_leaf=0.0,
                        random_state=1, splitter='best')
```

In [79]:


```
pd.DataFrame({'feature':feature_cols,
              'importance':churn_tree_clf.feature_importances_})
```

Out[79]:

	feature	importance
0	Account_Length	0.008279
1	VMail_Message	0.069224
2	Day_Mins	0.326225
3	Eve_Mins	0.190501
4	Night_Mins	0.031833
5	Intl_Mins	0.101552
6	CustServ_Calls	0.147903
7	Intl_Plan	0.000000
8	VMail_Plan	0.124483

In [80]:

```
y_predict = churn_tree_clf.predict( X_test )
cm = metrics.confusion_matrix( y_test, y_predict )
sn.heatmap(cm, annot=True, fmt='.2f' );
```



In [81]:

```
metrics.accuracy_score( y_test, y_predict )
```

Out[81]:

0.92500000000000004

## Create tree graph

Download and Install **Graphviz** from <http://www.graphviz.org/Download..php>  
(<http://www.graphviz.org/Download..php>)

Add the path (for example C:\Program Files (x86)\Graphviz2.38\bin) to your PATH variable.

In [82]:


```
from sklearn.tree import export_graphviz

export_graphviz(churn_tree_clf, out_file='churn_tree.odt',
                feature_names=feature_cols)
```

*Now convert the dot file to png file*

**dot -Tpng filename.dot -o outfile.png**

Now you can open and see the decision tree diagram in any editor.



## Random Forest Classifier

In [72]:

```
from sklearn.ensemble import RandomForestClassifier
rfClf = RandomForestClassifier( n_estimators= 50 )
scores = cross_val_score(rfClf,
                          X_train,
                          y_train,
                          cv=10,
                          scoring='accuracy')

scores
```

Out[72]:

```
array([ 0.92735043,  0.92307692,  0.94017094,  0.92735043,  0.9230769
2,
        0.9527897 ,  0.93562232,  0.92274678,  0.90517241,  0.9310344
8])
```

In [73]:

```
scores.mean()
```


Out[73]:

```
0.92883913358606274
```



In [74]:

```
rfClf.fit( X_train, y_train )
y_predict = rfClf.predict( X_test )
cm = metrics.confusion_matrix( y_test, y_predict )
sn.heatmap(cm, annot=True, fmt='.2f' );
```



In [85]:

```
metrics.accuracy_score( y_test, y_predict )
```

Out[85]:

0.92500000000000004

## Naive Bayes Model

In [76]:

```
from sklearn.naive_bayes import GaussianNB
nbClf = GaussianNB()
scores = cross_val_score(nbClf,
                          X_train,
                          y_train,
                          cv=10,
                          scoring='accuracy')

scores
```

Out[76]:

```
array([ 0.83760684,  0.84188034,  0.8034188 ,  0.81196581,  0.8547008
5,
        0.86266094,  0.84978541,  0.87982833,  0.87068966,  0.8362069
])
```

In [86]:

```
scores.mean()
```

Out[86]:

```
0.84487438794083736
```